

# WordCount in Spark

**2/4 points (50%)**

Quiz, 4 questions

## ✖ Try again once you are ready.

Required to pass: 80% or higher

You can retake this quiz up to 3 times every 8 hours.

[Back to Week 4](#)[Retake](#)

1. What does the following line of code do?

1 / 1  
points

```
words = lines.flatMap(lambda line: line.split(" "))
```

- ☐ Each word is merged into lines to be counted later.
- ☐ Each word in each line is counted.
- ☒ Each line in the document is split up into words.

Correct

- ☐ Each line in the document is split into various Spark partitions.



2. What does the following line of code imply about the state of partitions before the action is performed?

0 / 1  
points

```
words = lines.flatMap(lambda line: line.split(" "))
```

- ☐ There is only one single partition containing the full document.

- ☒ Each Spark partition corresponds to a word in the document.



**This should not be selected**

- ☐ Each Spark partition corresponds to a line in the document.



3. When the following command is executed, where is the file written and how can it be accessed?

0 / 1  
points

```
counts.coalesce(1).saveAsTextFile("hdfs://user/cloudera/wordcount/outputDir")
```

- ☐ HDFS and through the system directory with the "cd" terminal command.
- ☐ HDFS and through the "hadoop fs" command.
- ☐ The local file system and through the directory with the "cd" terminal command.
- ☒ The local file system and through the "hadoop fs" command.



**This should not be selected**



4. What does the number one (1) allow us to do in the following line of code?

1 / 1  
points

```
tuples = words.map(lambda word: (word,1))
```

- ☐ The number represents the number of partitions in charge of keeping track of each word.
- ☐ The number represents the number of partitions in charge of counting each line.
- ☐ None, completely arbitrary in order to apply an algorithm that requires a tuple.
- ☒ Treat each word with a weight of one during the counting process.

**Correct**

